

Physics4ML

Lucas Leung^a

^aRudolf Peierls Centre for Theoretical Physics, University of Oxford, Parks Road, Oxford OX1 3PU, UK

E-mail: lucas.leung@physics.ox.ac.uk

ABSTRACT: A short set of seminar notes on Physics4ML.

Contents

1	Basics of ML and NNs	1
1.1	NN basics	2
2	The three pillars of neural networks	5
2.1	Expressivity of NNs	6
2.1.1	Universal Approximation Theorem	6
2.1.2	Kolmogorov-Arnold Theorem	7
2.2	Statistics of NNs	7
2.2.1	NNGP correspondence	8
2.2.2	Non-Gaussian processes	11
2.2.3	Symmetries - a first encounter	12
2.3	Dynamics of NNs	13
2.3.1	Neural tangent kernel	14
2.3.2	Feature learning	15
3	NN-FT correspondence	16
3.1	Parametrically breaking Gaussianities	18
3.2	EFTs and Feynman diagrams for NN-FT	18
3.3	Engineering actions in NN-FT	18
3.4	Explicit examples	18
3.5	Classical field configurations and Landscapes	18
4	Conformal NNFTs	18

1 Basics of ML and NNs

We are interested in Machine Learning (ML). What is ML? Mitchell gives a succinct definition.

Definition 1.1. A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P ; if its performance at tasks in T , as measured by P , improves with experience E .

Bit of a weird definition. Let us look at what the terms mean.

- **The task T .** This is some kind of problem that we wish to solve. Of course, typically this will be too hard to solve with fixed programs written and designed by human beings. Classes of tasks include for example: classification, regression, transcription, etc.

- **The performance measure P .** This is a quantitative measure of the performance of a computer program, specific to the task being carried out by the system. Often times we measure the accuracy of a model, using a test set that the computer program has not been trained on.
- **The experience E .** ML algorithms are broadly categorised as **supervised** and **unsupervised**. A **dataset** is the set of data drawn from the world, collected for the purposes of the task T . Whilst unsupervised learning algorithms aim to experience the data set with many features and extract useful properties of the set, supervised learning algorithms often associate each data point with an additional label or target (see example below). The line between the two is not clear, and there are many others (like for example Reinforcement Learning where the algorithm interacts with an environment).

There are many additional issues that come with ML. Studying ML involves understanding generalisation and statistical learning theory. We want to minimise both the error coming from the data being trained on (**training error**) as well as its generalisation (**generalisation error**) - and the two extremes of training are called **underfitting** and **overfitting**. A choice of model, algorithm, hyperparameters (settings), regularisation etc. all contribute to the **representational capacity** of the model. Depending on the problem, a specific choice of algorithm might work a lot better than a different one, so optimising for an algorithm for a specific task is a big challenge in applications of machine learning to solve problems. Additionally, as Bayesian inference is often used in predicting models, a good understanding of maximum likelihood estimation is also needed in designing good ML algorithms.

There are however a lot of problems with basic ML:

- **Curse of dimensionality.** ML problems just become exceedingly difficult due to logarithmic scaling.
- **Local constancy prior is not sufficient.** Often times simply requiring a smooth function is not enough (c.f. classification problems - the number of contiguous regions do not grow with training samples).
- **Manifold structure.** Learning seems to be hopeless when focusing on the entire region - but focusing on a subspace ('manifold') seems to be the way to go. This is not well-understood in ML.

One direction of improvement is to understand **deep learning**, learning using deep neural networks. So we should talk about neural networks.

1.1 NN basics

Artificial neural networks are loosely modelled after the structures of neurones and synapses in our brains. Let's begin with a simple definition.

Definition 1.2. A **neural network** is a function

$$\phi_\theta : \mathbb{R}^d \rightarrow \mathbb{R}^D \quad (1.1)$$

with parameters θ . We can $x \in \mathbb{R}^d$ the **input**, $\phi_\theta(x) \in \mathbb{R}^D$ the **output**, $\phi_\theta \in \text{Maps}(\mathbb{R}^d, \mathbb{R}^D)$ the **network** and \mathcal{D} the **data**. The **data** \mathcal{D} depends on the problem and involves at least a subset of \mathbb{R}^d potentially paired with labels $y \in \mathbb{R}^D$.

The easiest case of course is when $D = 1$. This is then just a scalar field - and we can already learn a lot from scalar function.

The question now we want to ask is the following.

Question: What does a NN predict?

For a set of fixed parameters θ then the answer is clear - this is just given by the function $\phi_\theta(x)$. The true answer however is complicated by two aspects in particular:

Statistics. NNs need to be initialised. When they are initialised, the parameters θ are drawn from some distribution $P(\theta)$,

$$\theta \sim P(\theta) \quad (1.2)$$

Different draws will give different θ and hence different neural networks ϕ_θ . So the prediction ϕ_θ is not fundamental - what is fundamental is the moments,

$$\mathbb{E}[\phi_\theta(x_1) \dots \phi_\theta(x_n)] = \int d\theta P(\theta) \phi_\theta(x_1) \dots \phi_\theta(x_n) \quad (1.3)$$

Of course, the first and second moments are then just the average prediction and variance respectively. Expectations are across different initialisations, so we can see that moments are fundamental quantities in the statistical distribution. We can replace $\mathbb{E}[\cdot] = \langle \cdot \rangle$ later using our language - to indicate that the moments are just correlation functions of the statistical ensemble of neural networks.

Dynamics. Of course, the ML parameters θ are updated during the training process to solve problems. So we definitely have a trajectory in the **parameter space** $\theta(t) \in \mathbb{R}^{|\theta|}$ (and therefore in the output and function space). The flow, or trajectory, will be governed by some learning dynamics determined by the optimisation algorithm and the nature of the problem. An example of an optimisation algorithm is gradient descent. For example, given a problem in supervised learning where we have the data

$$\mathcal{D} = \left\{ (x_\alpha, y_\alpha) \in \mathbb{R}^d \times \mathbb{R}^D \right\}_{\alpha=1}^{|\mathcal{D}|}, \quad (1.4)$$

and a loss function

$$\mathcal{L}[\phi_\theta] = \sum_{\alpha=1}^{|\mathcal{D}|} l(\phi_\theta(x_\alpha), y_\alpha), \quad (1.5)$$

where l is the loss function (defined as, for example, the mean square error). The parameters θ can be optimised by gradient descent,

$$\frac{d\theta_i}{dt} = -\nabla_{\theta_i} \mathcal{L}[\phi_\theta]. \quad (1.6)$$

This defines a certain flow in the parameter space as defined by the loss function and the optimisation algorithm (the gradient descent).

Now of course, we can see that the statistics of the draw can be thought of as evolving under training as well. In particular, we can think of the parameters $\theta(t)$ as drawn from the time-dependent distribution $P(\theta(t))$,

$$\theta(t) \sim P(\theta(t)) , \quad (1.7)$$

where t is the training time parameter. This gives time-dependent correlators,

$$G_t^{(n)}(x) = \langle \phi_\theta(x_1) \dots \phi_\theta(x_n) \rangle_t , \quad (1.8)$$

with the subscript t indicating time-dependence and the expectation is with respect to the distribution $P(\theta(t))$. If learning helps, then the $t \rightarrow \infty$ limit should give interesting correlation functions - this is what we will focus on later and we will find that in a certain supervised setting there is an exact analytical solution for this quantity.

So far, we have only talked about initialisation and training of neural networks. There is another important aspect of neural networks that we have missed - **architecture**. In particular, the correlation functions that we have computed above, which we can now write as,

$$G^{(n)}(x_1, \dots, x_n) = \frac{1}{Z_\theta} \int d\theta \phi(x_1) \dots \phi(x_n) P(\theta) , \quad (1.9)$$

where the partition function is written as $Z_\theta = \int d\theta P(\theta)$. Where does the randomness of the fixed field configuration ϕ arise from? We can think of the functions ϕ as random functions (in field theory, this will be the fields themselves). The statistics of the ensemble of the ϕ here, however, unlike in the field theory where the probability density function is specified $P(\phi) = \exp(-S[\phi])$, is instead determined by the construction of the networks and the parameter space distribution. This means that the **architecture** of the neural network defines the functional from ϕ , as well as the choice of distributions from which the parameters θ are drawn from. This begs the question:

Question: How powerful is NN?

The answer to that of course depends on the **architecture**. We can have a simple NN where we only consider linear functional forms - this is linear so will fail to express any non-linearity in our problems. So this architecture, we say, is not **expressive** enough. We therefore have a third intricate pillar that ties in with statistics and dynamics of NNs - **Expressivity**.

What does a typical neural network architecture look like? The most common (and arguably, the simplest) implementation of NNs are called **feed-forward neural networks**. This consists of a set of layers, each of which consists of a set of nodes, whose number typically varies between the layers. Feed-forward NNs are directed - the information is passed on from the **input layer** to the **output layer** via intermediate **hidden layers**, known as such as the states in such layers are hidden. The number of layers L is the **depth**, and the

number of nodes per layer $N^{(i)}$ is the **width** - a NN with a $L \gg 1$ is called a deep NN, or a **multi-layer perceptron (MLP)**, or a **deep feedforward network**. The output of each layer is often set to have the following form,

$$l_{\mu}^{(i)} = \sigma^{(i)} \left(\sum_{\nu=1}^{n_{i-1}} w_{\mu\nu}^{(i)} l_{\nu}^{(i-1)} + b_{\mu}^{(i)} \right) \quad (1.10)$$

where $l_{\nu}^{(i-1)}$ are the output of the $(i-1)$ th layer, the matrix entries $w_{\mu\nu}^{(i)}$ are known as **weights**, and the vector $b_{\nu}^{(i)}$ **bias**. The **activation function** $\sigma^{(i)}$ is non-linear function which introduces non-linearity to the NN. This is often picked depending on the purpose of the NNs (as some activation functions are shown to help with convergence in particular problems). The typical activation functions include the identity, log-sigmoid, tanh or leaky ReLU (see sections below). The above then defines an example of a typical NN architecture.

To summarise, we have the following three questions we want to know from neural networks.

1. **Expressivity**. How powerful is NN?
2. **Statistics**. How are NN's initialised and what is the NN statistical ensemble?
3. **Dynamics**. How do NNs evolve in training?

The centre of this set of lectures is to explain how it is possible to understand these questions from a theoretical physicist's point of view, utilising the tools we have on our hand:

- Field Theory.
- Landscape Dynamics from loss functions of the parameter space.
- Symmetries of individual NNS and their ensembles.

As theoretical physicists, such a new perspective is not just helpful for us to understand how NNs work, but also to understand empirical results in machine learning where these results are barely explored. We are not thriving for a rigorous proof to all the results, but to first develop some mechanisms and toy models so we can later add in rigour. That is the main goal of the programme of **Physics4ML**.

2 The three pillars of neural networks

Let us now try and understand the three pillars of neural networks! A quick recap of the three pillars of the neural networks:

1. **Expressivity**. How powerful is NN?
2. **Statistics**. How are NN's initialised and what is the NN statistical ensemble?
3. **Dynamics**. How do NNs evolve in training?

Let us try and understand these one-by-one. This ties in closely with the recent results currently being developed in the computer science community. The expressivity of NNs is first explored by Cybenko using the Universal Approximation Theorem (UAT), where as the statistics and dynamics of NNs are recently covered by the Neural Network Gaussian Processes (NNGP) and Neural Tangent Kernel (NTK) programme. In this section, we will understand these results in these directions, with the help of physical intuition from physics.

2.1 Expressivity of NNs

We have alluded above in §1 that neural networks are big functions composed out of many simple functions given a choice of the architecture. A different architecture would give a different composition of simple functions and hence a different expression of the neural network. The question then becomes a mathematical one - how good is NN at approximating an unknown function?

2.1.1 Universal Approximation Theorem

The Universal Approximation Theorem, UAT for short, is the first result in showing how good are NNs at approximating unknown functions. It states that a NN with a single hidden layer can approximate any continuous function on a compact domain to arbitrary accuracy. The precise form of the theorem is stated by Cybenko:

Theorem 2.1 (Cybenko). *Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a continuous function on a compact set $K \subset \mathbb{R}^d$. Then for any positive number $\epsilon > 0$ there exists a neural network with a single hidden layer of the form,*

$$\phi(x) = \sum_{i=1}^N \sum_{j=1}^d w_i^{(1)} \sigma(w_{ij}^{(0)} x_j + b_i^{(0)}) + b^{(1)}, \quad (2.1)$$

where the parameters are $\theta = \{w_{ij}^{(0)}, w_i^{(1)}, b_i^{(0)}, b^{(1)}\}$, where $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ is a non-polynomial nonlinear function such that

$$\sup_{x \in K} |f(x) - \phi(x)| < \epsilon. \quad (2.2)$$

Proof. See [1]. The idea of the proof is the following. In the original work, Cybenko focused on the case where σ is the sigmoid function,

$$\sigma(x) = \frac{1}{1 + e^{-x}}, \quad (2.3)$$

which appear in the network function in the following form,

$$\sigma(w_i^{(0)} \cdot x + b_i^{(0)}) \quad (2.4)$$

and this approximates a shifted step function as $w^{(0)} \rightarrow \infty$. So scaling this with $w^{(1)}$, this effectively gives a bump - which can be put together to approximate any function (which is very similar to the Riemann integral). \square

Note that the UAT is a powerful result, but it has some big limitations. In particular, firstly, although the error ϵ improves with N the precise accuracy bound is unknown. Secondly, it tells us there is a point θ^* in parameter space that is a good approximation to any f but it doesn't tell you how to get there - we still need to answer the questions about learning dynamics.

However, this is still a first result. In particular, there are many generations of the theorem to other activations, deep NNs, and other domains.

2.1.2 Kolmogorov-Arnold Theorem

Are there other architectures that we can construct - and are there any related theorems that support such construction, if it exists? One such construction is due to Kolmogorov and Arnold - any multivariate continuous function can be represented exactly as a sum of continuous one-dimensional functions.

Theorem 2.2 (Kolmogorov-Arnold Representation Theorem). *Let $f : [0, 1]^n \rightarrow \mathbb{R}$ be an arbitrary multivariate continuous function. Then it has the representation*

$$f(x_1, \dots, x_n) = \sum_{q=0}^{2n} \Phi_q \left(\sum_{p=1}^n \phi_{q,p}(x_p) \right) \quad (2.5)$$

where Φ_q and $\phi_{q,p}$ are continuous one-dimensional functions.

A non-trivial example of the KAT is the function $f(x, y) = xy$. Since,

$$f(x, y) = xy = e^{\log x + \log y}, \quad (2.6)$$

so the theorem works for multiplications. Since this is a sum of functions, one can try and build a neural network with this theorem. However, note that $\phi_{p,q}$ lives on the connection between x_p and $x_q^{(1)} = \sum_{p=1}^n \phi_{q,p}(x_p)$. Before, when we had the feedforward network, we can think of the linear transformations (the weights and biases) as acting on the edges and the activation function acting on the nodes of each of the layers. Here, instead the activation functions $\phi_{q,p}$ act on the edges instead of the nodes. For distinct values of p, q the functions $\phi_{q,p}$ are independent in general as well. Recently, this architecture is explored in the proposal of **Kolmogorov-Arnold networks** (KAN) [2]. This has the advantage of interpretable architectures which can be mapped directly into a symbolic representation, improving the interpretation of the neural network.

2.2 Statistics of NNs

Now we move on to understanding what happens to the neural network at initialisation. In particular, we are interested in the statistical properties of NNs. We want to in particular answer the following question:

Question: What characterises the statistics of the NN ensemble?

One particular aspect of this is encoded in the moments or n -point functions, obtained from the partition functions as the following (using field theory notions):

$$Z[J] = \left\langle e^{\int d^d x J(x) \phi(x)} \right\rangle, \quad (2.7)$$

$$G^{(n)}(x_1, \dots, x_n) = \left(\frac{\delta}{\delta J(x_1)} \cdots \frac{\delta}{\delta J(x_n)} Z[J] \right) \Big|_{J=0} = \langle \phi(x_1) \dots \phi(x_n) \rangle, \quad (2.8)$$

where, using usual field theory notations, $J(x)$ would be the source. This expectation, as we have discussed before, is actually here averaged across different distributions $P(\theta)$ of the parameters θ , giving the partition function as

$$Z[J] = \int d\theta P(\theta) e^{\int d^d x J(x) \phi(x)}. \quad (2.9)$$

However, we are more used to functional (Feymann path) integrals of the form

$$Z[J] = \int \mathcal{D}\phi e^{-S[\phi]} e^{\int d^d x J(x) \phi(x)}. \quad (2.10)$$

where the action $S[\phi]$ defines a density on functions. We call the first the **parameter space partition function**. The question now is:

Question: What is the action $S[\phi]$ associated to a distribution $(\phi, P(\theta))$?

To answer this, we actually must first take a step back and understand something known as the Neural Network-Gaussian Process (NNGP) correspondence. From this, we will see that the parameter-space and function-space descriptions can be thought of as a duality effectively.

2.2.1 NNGP correspondence

In this section the aim is to discuss the correspondence between Neural Networks and Gaussian Processes. This is first pointed out by Neal [3] in his thesis, where he had the following simple observation of the a single-layer fully connected network of width N :

$$\phi(x) = \frac{1}{\sqrt{N}} \sum_{i=1}^N \sum_{j=1}^d w_i^{(1)} \sigma(w_{ij}^{(0)} x_j) \quad (2.11)$$

where for simplicity I have set the biases $b_j^{(0)}$ to zero. The set of network parameters $\theta = \{w_{ij}^{(0)}, w_i^{(1)}\}$ is independently and identically distributed. We can now use the functional form of the Central Limit Theorem (CLT) where since $\phi(x)$ is the sum of random functions that are independently and identically distributed, in the limit where the network parameters $N \rightarrow \infty$ we then have that $\phi(x)$ is a Gaussian. Let us understand these terms in more detail.

Gaussian Process. A time continuous stochastic process with variables $\{X_t\}_{t \in T}$ is a Gaussian process if and only if for every finite set of indices in the index set T , $(X_{t_1}, \dots, X_{t_k})$ is a multivariate Gaussian random variable. Explicitly, recall the Bayesian approach to

model physical theories, where we try and use a nonlinear function $f(x)$ parameterised by the parameters w_a to model the data (x_i, y_j) . By Bayes' theorem, the inference of the function $f(x)$ is described by

$$P(f(x)|x_i, y_j) = \frac{P(y_j|f(x), x_i) P(f(x))}{P(y_j|x_i)} . \quad (2.12)$$

The idea of Gaussian process modelling is to place a prior $P(f(x))$ directly on the space of functions, without parameterising x . Loosely, Gaussian process is the simplest type of prior over functions where it is a generalisation of a Gaussian distribution over a finite vector space to a function space of infinite space. A Gaussian process is specified by a **mean** function $\mu(x)$ and **covariance matrix** $C(x, x')$, also known as a **kernel**. Being even a bit more mathematically imprecise, we can write a Gaussian process as the following probability distribution on the space of functions $f(x)$:

$$P(y(x)|\mu(x), A) = \frac{1}{Z} \exp \left[-\frac{1}{2} (y(x) - \mu(x))^T A (y(x) - \mu(x)) \right] , \quad (2.13)$$

where $\mu(x)$ is the mean function and A is a linear operator and the inner product of two functions is defined by

$$y(x)^T z(x) = \int dx y(x) z(x) . \quad (2.14)$$

We can in fact write

$$\log P(y(x)|\alpha) = -\frac{1}{2} y(x)^T A y(x) + \text{const} , \quad (2.15)$$

where $A = [D^p]^T D^p$ and D is the linear derivative operator.

NNGP Correspondence. Now that we know what is a Gaussian process is, we will now give the NNGP correspondence. This is as follows. Given a NN of single-layer, fully-connected of width N , in the limit $N \rightarrow \infty$, we have a Gaussian process:

$$\lim_{N \rightarrow \infty} \phi(x) \sim \mathcal{N}(\mu(x), K(x, y)) , \quad (2.16)$$

with mean and covariance $\mu(x)$ and $K(x, y)$.

Explanation. The NNGP Correspondence can be thought of as the following. For many architectures, for example,

- Single layer fully-connected NN: $N = \text{width}$
- Deep fully connected network: $N = \text{width}$
- Convolutional NN: $N = \text{channels}$
- Attention networks: $N = \text{heads}$

where the limit $N \rightarrow \infty$ exists, the distribution on the functions becomes a Gaussian process. In particular, the NN outputs $\{f(x_1), \dots, f(x_k)\}$ evaluated on any fixed set of k inputs x_i are drawn from a multivariate Gaussian distribution $\mathcal{N}(\mu, \Xi^{-1})$,

$$\{f(x_1), \dots, f(x_k)\} \sim \mathcal{N}(\mu, \Xi^{-1}) . \quad (2.17)$$

The inverse covariance matrix Ξ is determined by the kernel function $K(x, x')$. If $\mu = 0$, then the GP is entirely determined by the covariance, so it is determined entirely by the kernel. To compute correlation functions, we can see that between n outputs this is expressed as,

$$G^{(n)}(x_1, \dots, x_n) = \frac{\int df f_1 \dots f_n e^{-\frac{1}{2} f_i \Xi_{ij} f_j}}{Z} \quad (2.18)$$

with the partition function being $Z = \int df e^{-S}$ and $S = -\frac{1}{2} f_i \Xi_{ij} f_j$ being the **log-likelihood** (c.f. the action). We can of course take the continuum limit to get

$$G^{(n)}(x_1, \dots, x_n) = \frac{\int df f_1 \dots f_n e^{-S}}{Z} \quad (2.19)$$

with the log-likelihood as (d is dimension of input space)

$$S = \frac{1}{2} \int d^d x d^d x' f(x) \Xi(x, x') f(x') \quad (2.20)$$

and $\Xi(x, x') = K^{-1}(x, x')$ being the inverse covariance function defined using the D -dimensional Dirac delta function:

$$\int d^d x' K(x, x') \Xi(x', x'') = \delta^{(d)}(x - x'') . \quad (2.21)$$

Let us see this explicitly worked out for the two-point correlator. We see that

$$G^{(2)}(x, y) = \frac{\Xi_2}{N} \left\langle \sigma(w_{ij}^{(0)} x_j) \sigma(w_{il}^{(0)} x_l) \right\rangle = \Xi_2 \left\langle \sigma(w_{ij}^{(0)} x_j) \sigma(w_{il}^{(0)} x_l) \right\rangle \quad (2.22)$$

where we have evaluated the sum over i in the last step (so there is no summation over i in the final expression). We can of course evaluate the quantity inside the bracket to get the answers (we can always numerically obtain an answer by, for example, a Monte Carlo estimate [4]).

Relation to Field Theory. So how does all of this relate to physics? Note that the correlators are defined exactly like the ones we have in field theories in physics, apart from the fact that in physics the S is the action that is defined at the start. If we combine the parameter distribution $\theta \sim P(\theta)$ and the network architecture together this induce an implicit distribution on the function space from which the neural network is drawn from, say $\phi \sim P(\phi)$. What that means is the following. Since the NNGP Correspondence tells us that $\exp(-S[\phi])$ is Gaussian in the $N \rightarrow \infty$ limit this means that $S[\phi]$ is quadratic in the networks, and in particular we can define the associated action using the two-point correlator as

$$S[\phi] = \int d^D x d^D y \phi(x) G^{(2)}(x, y)^{-1} \phi(y) , \quad (2.23)$$

and the NNGP correspondence gives

$$\lim_{N \rightarrow \infty} \phi(x) \sim \mathcal{N}\left(0, G^{(2)}(x, y)\right) . \quad (2.24)$$

So in the $N \rightarrow \infty$ limit, certain large NNs are just functions drawn from generalised free-field theories. An example, of couses, is the free scalar field theory given by

$$Z = \int D\phi e^{-S[\phi]} \quad (2.25)$$

where the action S is defined as

$$S[\phi] = \int d^d x \phi(x) (\square + m^2) \phi(x) \quad (2.26)$$

2.2.2 Non-Gaussian processes

It is obvious that the $N \rightarrow \infty$ limit, in the most practical sense, does not hold. Violating any of the assumptions of the CLT should introduce non-Gaussian contributions, i.e. in the field theory language, interactions. In particular, the CLT is violated by

- $\frac{1}{N}$ -corrections.
- Independence breaking - in the CLT proof the cumulant generating function is assumed to be the sum of all cumulant generating functions which ignores cross-correlations.

So we, in general, expect non-Gaussianities. Non-Gaussianities can be computed via the connected four-point function and using effective field theory. In particular perturbative diagrammatic methods have been used to study NNs at initialisation. Let us in the following highlight at least one of these points - the connected four-point function. The full derivation is in [5], but the result is, given an architecture,

$$\phi(x) = \sum_i w_i \varphi_i(x) , \quad (2.27)$$

$$G_c^{(4)}(x, y, z, w) = \frac{1}{N} \left[\Xi_4 \langle \varphi_i(x) \varphi_i(y) \varphi_i(z) \varphi_i(w) \rangle - \Xi_2^2 (\langle \varphi_i(x) \varphi_i(y) \rangle \langle \varphi_i(z) \varphi_i(w) \rangle + \text{perms}) \right] . \quad (2.28)$$

Note that this is non-zero at finite- N - so there are interactions.

When non-Gaussianities are small, the theory will be close to free theory and be weakly interacting, in which case we can proceed to calculate the correlation functions using Feynman diagrams. The higher connected correlation functions which vanish in the Gaussian limit now capture non-Gaussianities - they are known as **cumulants** and can be obtained from generating functions $W[J]$ as

$$G_c^{(n)}(x_1, \dots, x_n) := \left(\frac{\delta}{\delta J(x_1)} \cdots \frac{\delta}{\delta J(x_n)} W[J] \right) \Big|_{J=0} \quad (2.29)$$

For a theory with a known Lagrangian description, the correlators $G_c^{(n)} \neq 0$ for $n > 2$ still encode the presence of non-Gaussianities. For NN-FTs, since we know the parameter space description always exists, we can study non-Gaussianities using connected correlators. This allows us to:

- NN \rightarrow FT: understand interactions in associated field theories.
- FT \rightarrow NN: capture the statistics of finite networks and networks with correlations in the parameter descriptions which generically develop during training.

We will come back to this when we discuss the details of the NN-FT correspondence in §3.

2.2.3 Symmetries - a first encounter

Now we can ask the question.

Question: Are there any structures in the ensemble?

One type of structure we can see here at level of initialisation is the symmetry of the architecture. In particular, recall that the networks $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^D$ are functions. It is perhaps possible to introduce some sort of symmetry on the NN architecture such that the family of NNs generated will be invariant under such symmetry. Or that the symmetry can be realised by the function. To do this, we need to use the notion of equivariance in the mathematical literature.

Equivariant NNs. We say that ϕ is **G -equivariant** with respect to a group G if

$$\rho_D(g)\phi(x) = \phi(\rho_d(g)x) , \quad (2.30)$$

for all $g \in G$ where $\rho_d \in \text{Mat}(\mathbb{R}^d)$ and $\rho_D \in \text{Mat}(\mathbb{R}^D)$ are matrix reps of G in \mathbb{R}^D and \mathbb{R}^d respectively. The network is then **invariant** if $\rho_D = \mathbb{1}$, i.e. if the rep is trivial.

Global symmetries in NN ensembles. Now let us look at the symmetries that arise in ensembles of NNs. In particular, we want to leave the statistical ensemble of NNs invariant. This is equivalent to the notion of **global symmetries** in field theory. Suppose the network transform under a group action as

$$\phi \mapsto \phi_g, \quad g \in G . \quad (2.31)$$

The ensemble of networks then has a global symmetry group G if the partition function is invariant,

$$Z_g[J] = Z[J] , \forall g \in G , \quad (2.32)$$

or we can write this as

$$\left\langle \exp \left(\int d^d x J(x) \phi_g(x) \right) \right\rangle = \left\langle \exp \left(\int d^d x J(x) \phi(x) \right) \right\rangle , \forall g \in G . \quad (2.33)$$

We can see that we can redefine the network and put the symmetry if $\langle \cdot \rangle$ is invariant. In terms of FT, this means that the action $S[\phi]$ and the measure $\mathcal{D}\phi$ are both invariant. In particular, it is possible to redefine the network by transforming the parameters, as follows:

$$\int d\theta_g P(\theta_g) \left\langle \exp \left(\int d^d x J(x) \phi_{\theta_g}(x) \right) \right\rangle = \int d\theta P(\theta) \left\langle \exp \left(\int d^d x J(x) \phi_{\theta}(x) \right) \right\rangle , \quad (2.34)$$

so the symmetry arises with the redefinition $\theta \mapsto \theta_g$ when the parameter density and measure must be invariant.

It is perhaps important to note that the formalism above allows for transforming the inputs and outputs of the network - which corresponds to internal and spacetime symmetries of the field theory. The notion of equivariance plays a central role here - the ensemble of equivariant NNs is invariant under ρ_d on the input if the partition function is invariant under the induced ρ_D on the output.

Example 2.1. Let us perhaps look at an example with the NN architecture being of the form:

$$\phi(x) = \sum_i w_i \varphi_i(x) \quad (2.35)$$

with w_i being $P(w)$ even. The \mathbb{Z}_2 -action $\phi \mapsto -\phi$ then may be absorbed into the parameters $w_g = -w_i$ with $dwP(w)$ invariant by the evenness, so these NNs only have even correlators.

There are many more examples in [5].

2.3 Dynamics of NNs

Now we move on to the next aspect of NNs. What happens to NNs during training? As explained §1, it is clear that NNs evolve in some way under gradient descent. This is often simplified in a scheme known as Neural Tangent Kernel (NTK) where we will then use to solve a model exactly in the case of MSE loss. In particular, we will focus on the problem of dynamics of supervised learning with gradient descent, with data

$$\mathcal{D} = \{(x_\alpha, y_\alpha)\}_{\alpha=1}^{|\mathcal{D}|} , \quad (2.36)$$

and the loss function

$$\mathcal{L}[\phi] = \frac{1}{|\mathcal{D}|} \sum_{\alpha=1}^{|\mathcal{D}|} l(\phi(x_\alpha), y_\alpha) . \quad (2.37)$$

We optimise the network parameters θ by a gradient descent procedure,

$$\frac{d\theta_i}{dt} = -\eta \nabla_{\theta_i} \mathcal{L}[\phi] , \quad (2.38)$$

where let us define the natural object of gradient descent in function space as

$$\Delta(x) = -\frac{\delta l(\phi(x), y)}{\delta \phi(x)} \quad (2.39)$$

which kind of gives the gradient of the loss in space of learned outputs. This yields,

$$\frac{d\theta_i}{dt} = \frac{\eta}{|\mathcal{D}|} \sum_{\alpha=1}^{|\mathcal{D}|} \Delta(x_\alpha) \frac{\partial \phi(x_\alpha)}{\partial \theta_i} . \quad (2.40)$$

Could we possibly use this equation to derive a quantity to simplify our analysis?

2.3.1 Neural tangent kernel

Let us define a quantity in ML called the **Neural Tangent Kernel (NTK)**. By training the network by gradient descent and using Eq.(2.40), we have,

$$\frac{d\phi(x)}{dt} = \frac{\partial\phi(x)}{\partial\theta_i} \frac{d\theta_i}{dt} = \frac{\eta}{|\mathcal{D}|} \sum_{\alpha=1}^{|\mathcal{D}|} \Delta(x_\alpha) \Theta(x, x_\alpha) , \quad (2.41)$$

where the NTK is

$$\Theta(x, x_\alpha) = \frac{\partial\phi(x)}{\partial\theta_i} \frac{\partial\phi(x_\alpha)}{\partial\theta_i} . \quad (2.42)$$

Recall what a kernel is. We can define a kernel as a map $V \times V \rightarrow \text{Sym}(V \otimes V)$ where some of the dimensions can be suppressed. This maps the pair (x, x') to an $n_L \times n_L$ symmetric matrix $K(x, x')$. The dual space of V is spanned by the set of linear forms $\mu : V \rightarrow \mathbb{R}$ where $\mu = \langle v, \cdot \rangle$. So we can define a map $\Phi_K : V^* \rightarrow V$ as

$$f_\mu(\cdot) = \Phi_K(\mu(x)) = \mu K(x, \cdot) = \langle d, K(x, \cdot) \rangle . \quad (2.43)$$

where \cdot is now the variable. Now defining the **cost function** C as a functional $C : V \rightarrow \mathbb{R}$ and therefore the **kernel gradient** which is defined as

$$\nabla_K C|_{f_0} = \Phi_K \left(\partial_f C|_{f_0} \right) , \quad (2.44)$$

which we can write,

$$\nabla_K C|_{f_0}(x) = \frac{1}{N} \sum_{j=1}^N K(x, x_j) d|_{f_0}(x_j) \quad (2.45)$$

and we say a time-dependent function $f(t)$ follows the **kernel gradient descent w.r.t** K if it satisfies,

$$\partial_t f(t) = -\nabla_K C|_{f(t)} . \quad (2.46)$$

The NTK is a fundamental object. But it is physically terrible to work with. Why?

- Parameter-dependence. We need to sum over billions of parameters.
- Time-dependent. The NTK time evolves.
- Stochastic. The NTK inherits the statistical uncertainty at initialisation.
- Non-local. It communicates information between x and x_α (in particular, the correlation between the gradients of the function $\phi(x)$ and $\phi(x_\alpha)$).

All this makes NTK unwieldy. But there is a fix - the NTK simplifies in the $N \rightarrow \infty$ limit, where we call the lazy regime $|\theta(t) - \theta(0)| \ll 1$ as the number of parameters is large but the evolution keeps them in a local neighbourhood. The NN in this case is approximately a linear-in-parameters model where

$$\lim_{N \rightarrow \infty} \phi(x) \cong \phi_{\theta_0}(x) + (\theta - \theta_0)_i \left. \frac{\partial\phi(x)}{\partial\theta_i} \right|_{\theta_0} , \quad (2.47)$$

so then

$$\lim_{N \rightarrow \infty} \Theta(x, x') \cong \Theta(x, x')|_{\theta_0} \cong \beta_\theta(x, x') , \quad (2.48)$$

where we have approximated the last term by an expectation value and hence the network dynamics are now governed by the frozen NTK $\bar{\Theta}$,

$$\frac{d\phi(x)}{dt} = -\frac{\eta}{|\mathcal{D}|} \sum_{\alpha=1}^{|\mathcal{D}|} \frac{\delta l(\phi(x_\alpha), y_\alpha)}{\delta \phi(x_\alpha)} \bar{\Theta}(x, x_\alpha) . \quad (2.49)$$

where I have defined the frozen NTK as

$$\bar{\Theta}(x, x') = \beta_\theta(x, x') . \quad (2.50)$$

This is not good though. What Eq. (2.49) means for evolution is this - since $\bar{\Theta}$ is a fixed function that can be calculated exactly (it is deterministic), the NN effectively has zero parameters. So the NN actually does not learn any features in the hidden dimensions and therefore there is no non-trivial evolution that would cause the NTK to evolve. So we have, effectively, learnt nothing. An exact solvable model for frozen-NTK dynamics, in particular, can be found in [5] with MSE loss, where the converged network is

$$\phi_\infty(x) = \phi_0(x) + \bar{\Theta}(x, x_\alpha) \bar{\Theta}(x_\alpha, x_\beta)^{-1} (y_\beta - \phi_0(x_\beta)) , \quad (2.51)$$

which is known as kernel regression.

However, all is not lost. In particular, there are two key results that we can infer from here.

1. The NTK, in the limit $N \rightarrow \infty$, converges to an explicit deterministic limit. The existence of such is merely shown but not proven by the example we have above (see [6]), but this explicitly connects NTK with NNGP.
2. The NTK stays asymptotically constant during training in the same limit. This means that the frozen NTK behaviour is in fact universal in the limit, and as the intermediate layer pre-activations $z_\alpha^{(l)}$ are also Gaussian distributed, the statistics of a randomly initialised NN is described by a sequence of generalised free field theories where correlations are propagated down the network according to a recursion relation.

We see that the dynamics of this family of NNs (in this $N \rightarrow \infty$ limit) is consistent with the behaviour of free field theories. The initialisation is akin to the initial set-up, and the evolution of free theories are deterministic and described by a kernel. The question now becomes, what happens when we move away from the limit?

2.3.2 Feature learning

To extend to a general study of learning dynamics of NNs, we need to specifically engineer tractable properties so we can do wise N -scaling/expansions to understand the behaviour. In particular let us start with three properties:

1. Finite initialisation pre-activations. This means that the pre-activations $z^{(l)}(x_\alpha)$ in each layer, the inputs to the activation functions, will be set to $O_N(1)$.
2. Learning in finite time. We want to train the NN in a reasonable time-frame, $\frac{d\phi(x)}{dt} \sim O_N(1)$.
3. Feature learning in finite time. This is the same as above but we want to explicitly learn the features of the system, i.e. on the pre-activations $\frac{dz^{(l)}}{dt} \sim O_N(1)$.

From [5] and its cited works, the key point is that the constraints listed above have a one-parameter family of solutions which is completely fixed when the learning rate is additionally assumed. In particular, it is possible to undergo a detailed N -scaling analysis to see how we can lead to richer learning regimes known as dynamical mean field theory or the maximal update parametrisation.

3 NN-FT correspondence

Neural networks, as discussed above, generally have \mathbb{R}^n as their domain and as a result live in a space with Euclidean signature. This means they define statistical field theories (in the first instance) that may or may not have analytic continuations to quantum field theories in the Lorentzian signature. But it is clear that there is a direct correspondence between NNs and field theories [7].

The question that us theoretical physicist are most concerned about, however, is when neural architecture defines a quantum field theory (QFT), i.e. does there exist an analytic continuation to Lorentzian signature that defines a QFT. The key is the Osterwalder-Schrader (OS) theorem of axiomatic field theory which gives the necessary and sufficient conditions, expressed in terms of the correlators, for the existence of a QFT after continuation. The axioms include the following:

- **Euclidean Invariance.** Correlation functions must be invariant in Euclidean signature. This ensures it is Lorentz-invariant after analytic continuation.
- **Permutation Symmetry.** Correlation functions must be invariant under permutation of their arguments (collection of points in Euclidean space). Of course, this is automatic with scalar outputs.
- **Reflection Positivity.** Correlation functions must satisfy a positivity condition known as reflection positivity. This is necessary for unitarity and the absence of negative-norm states in the analytically-continued theory. For the Gaussian case, this simply requires the two-point function to be reflection positive, which means requiring

$$\int d^d x d^d y f^*(x) f(y) G^{(2)}(x^\theta, y) \geq 0. \quad (3.1)$$

for any complex function $f(x)$ with support only on $\tau > 0$ and $x^\theta = (-\tau_x, \mathbf{x})$ being the reflection of x in imaginary time.

- **Cluster Decomposition.** Correlation functions must satisfy cluster decomposition, which states that interactions must shut off at infinite distance. For connected correlators, this imposes,

$$\lim_{b \rightarrow \infty} G_c^{(n)}(x_1, \dots, x_p, x_{p+1} + b, \dots, x_n + b) \rightarrow 0, \quad (3.2)$$

for any value of $1 < p < n$.

In particular, our field theory only consists of fields and correlation functions at this stage, and is, really, an ensemble of functions with a way to compute their correlators. In the Euclidean picture, this is, additionally, a statistical ensemble of functions. We can already construct the partition function,

$$Z[J] = \left\langle \exp \left(\int d^d x J(x) \phi(x) \right) \right\rangle, \quad (3.3)$$

where we can use to compute correlators. The NN-GP correspondence now gives us the following correspondence of partition functions:

$$\int \mathcal{D}\phi e^{-S[\phi] + \int d^d x J(x) \phi(x)} \longleftrightarrow \int d\theta P(\theta) e^{\int d^d x J(x) \phi(\theta)}, \quad (3.4)$$

where we can now try and work out the associated action given $(\phi_\theta, P(\theta))$. Now I hear you complain - well, surely, we need more things! The list of things X may include,

- Quantumness
- Lagrangian Description
- Symmetries
- Locality

But these are all add-ons remember. Different physicists might argue that different things are important, and will hence (basically) require a different NN architecture and statistical ensemble to make $FT + X$ work. But FT is enough here. For now.

So now we know that there is an NN-FT correspondence, and there are additional things that different physicists might want in NNs to add extra spice to their field theories. Fine. We have constructed free field theories above, so the question is how could we possibly (perhaps the easiest way) get non-Gaussianities? Turns out it is possible to insert an operator to any local potential $V(\phi)$ which deforms the action in the following way:

$$Z[J] = \int d\theta P(\theta) e^{\int d^d x V(\phi_\theta(x))} e^{\int d^d x J(x) \phi_\theta(x)} \quad (3.5)$$

$$= \int d\theta \tilde{P}(\theta) e^{\int d^d x J(x) \phi_\theta(x)} \quad (3.6)$$

where we have used the architecture equation to define a new parameter density in the second line. The interactions in $V(\phi)$ will now break the Gaussianity of the NNGP as the statistical independence is now violated - giving the $Z[J]$ an interacting NN-QFT structure!

Below we will approach this correspondence in the following perspective.

3.1 Parametrically breaking Gaussianities

3.2 EFTs and Feynman diagrams for NN-FT

3.3 Engineering actions in NN-FT

3.4 Explicit examples

3.5 Classical field configurations and Landscapes

Spontaneous Symmetry Breaking (SSB).

A typical example is the ϕ^4 operator.

4 Conformal NNFTs

This is based on [8].

References

- [1] G. Cybenko, “Approximation by superpositions of a sigmoidal function,” *Mathematics of Control, Signals, and Systems (MCSS)*, vol. 2, pp. 303–314, Dec. 1989.
- [2] Z. Liu, Y. Wang, S. Vaidya, F. Ruehle, J. Halverson, M. Soljačić, T. Y. Hou, and M. Tegmark, “KAN: Kolmogorov-Arnold Networks,” 4 2024.
- [3] R. M. Neal, *Bayesian Learning for Neural Networks*. Springer New York, 1996 ed., 1996.
- [4] D. J. C. MacKay, *Information Theory, Inference, and Learning Algorithms*. Copyright Cambridge University Press, 2003.
- [5] J. Halverson, “TASI Lectures on Physics for Machine Learning,” 7 2024.
- [6] A. Jacot, F. Gabriel, and C. Hongler, “Neural tangent kernel: Convergence and generalization in neural networks,” in *Advances in Neural Information Processing Systems* (S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, eds.), vol. 31, Curran Associates, Inc., 2018.
- [7] M. Demirtas, J. Halverson, A. Maiti, M. D. Schwartz, and K. Stoner, “Neural network field theories: non-Gaussianity, actions, and locality,” *Mach. Learn. Sci. Tech.*, vol. 5, no. 1, p. 015002, 2024.
- [8] J. Halverson, J. Naskar, and J. Tian, “Conformal Fields from Neural Networks,” 9 2024.